

What is Chef..

..and how we use it at

tripsta[®]

Who am I?

Giedrius Rimkus

Lithuanian
PHP Developer at tripsta
Ruby enthusiast
Basketball lover

twitter  **@giedriusr**

I'll be talking about..

Sentiments

Pain

Solution

Back in 2009..

Infrastructure






Infrastructure




Infrastructure




Infrastructure



Infrastructure



Infrastructure



Complexity



Pain

- Long installation process
- Hard maintenance
- Scaling issues
- Constantly increased load and usage of resources (no load balancer)

Solution



CLUSTER

WTH IS CLUSTER?

What is a Cluster?

It's a group of linked computers, working together closely thus in many respects forming a single computer.

Cluster categories

High-availability (HA) clusters

High-availability clusters (also known as failover cluster) are implemented primarily for the purpose of improving the availability of services that the cluster provides.

Load-balancing clusters

Load-balancing is when multiple computers are linked together to share computational workload or function as a single virtual computer. Logically, from the user side, they are multiple machines, but function as a single virtual machine.

Compute clusters

Sometimes called as a “Grid computing”. Basically used for calculating huge stats, etc.

HOW TO BUILD IT?

**IN CASE OF
EMERGENCY**

Call **911**

OR



**I will ask for
help before my
problems get out
of control.**

RAGNARSON
PERFECT PROGRAMMERS

Yandy.com







Chef

What is Chef? What problem does it solve?

Chef is an open-source systems integration framework built specifically for automating the cloud / system configuration.

Chef types

Chef Solo

Chef client and Chef server

Hosted Chef

Private Chef

Chef Solo..

..is an open source standalone version of Chef that runs locally on your node, detached from a Chef server.

Chef Client and Chef Server

Chef-client connects to a Chef Server to be told what to do on the node.

Hosted Chef

As with Chef-Server, Chef-client connects to Hosted Chef to be told what to do on the local node.

Private Chef..

..is for Enterprises who want the power, flexibility, availability, and performance of Hosted Chef, but require that information never leave their private networks.

Why it's an issue?

Infrastructure changes all the time.

Different operating systems

Different hardware from different vendors.

What exactly can you do with Chef?

- Install Operating Systems on new servers.
- Install application software on servers.
- Have new software automatically configure itself to match your environment.
- Share recipes (and obtain recipes from) other people to install and configure software.

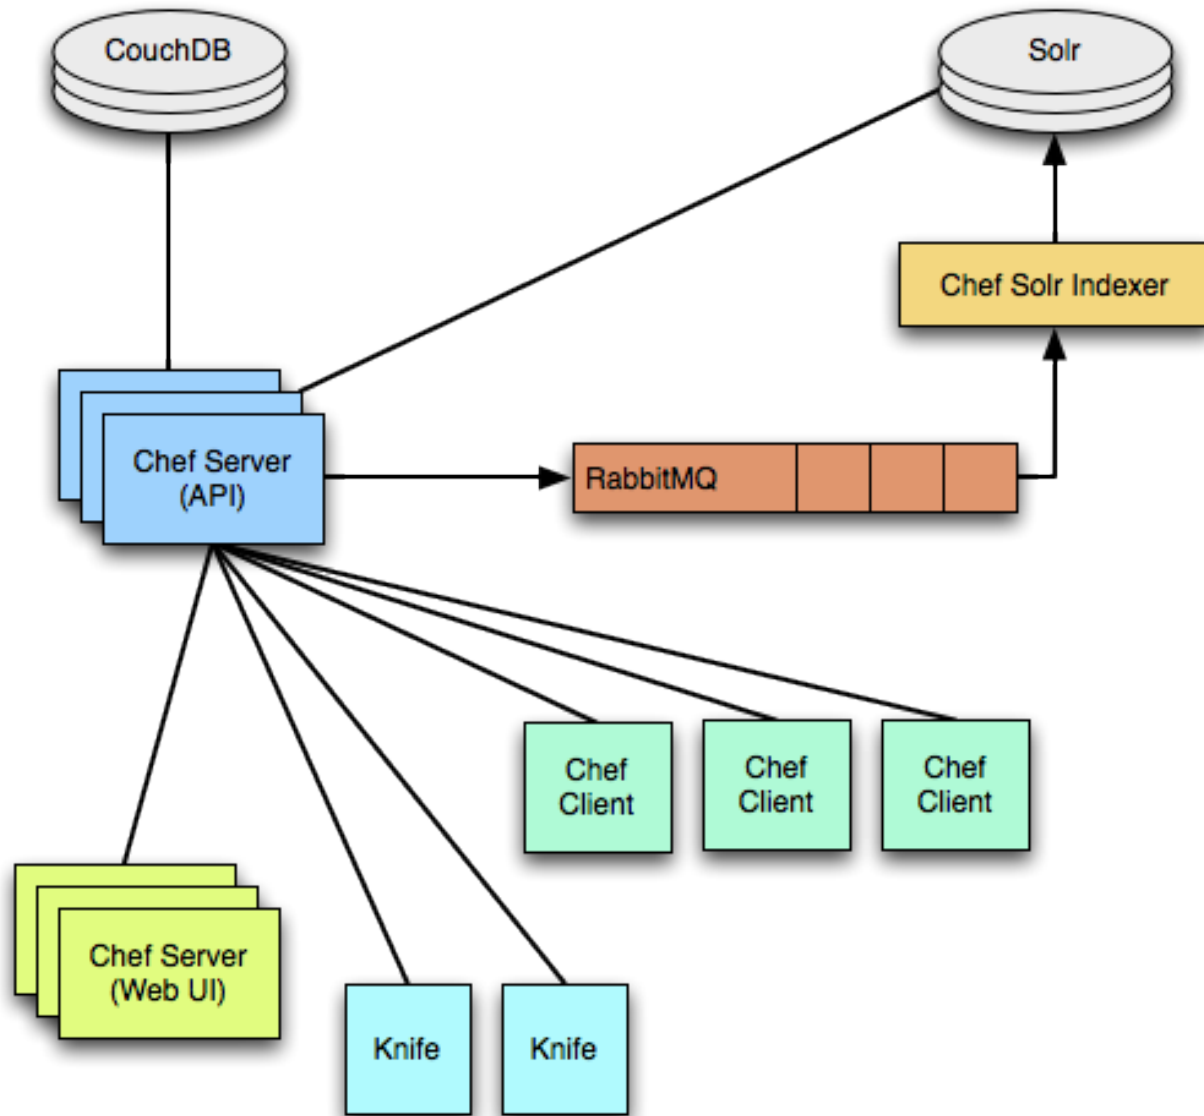
Some goodies I like about Chef

"Manage your servers by writing code,
not by running commands."

Chef is idempotent

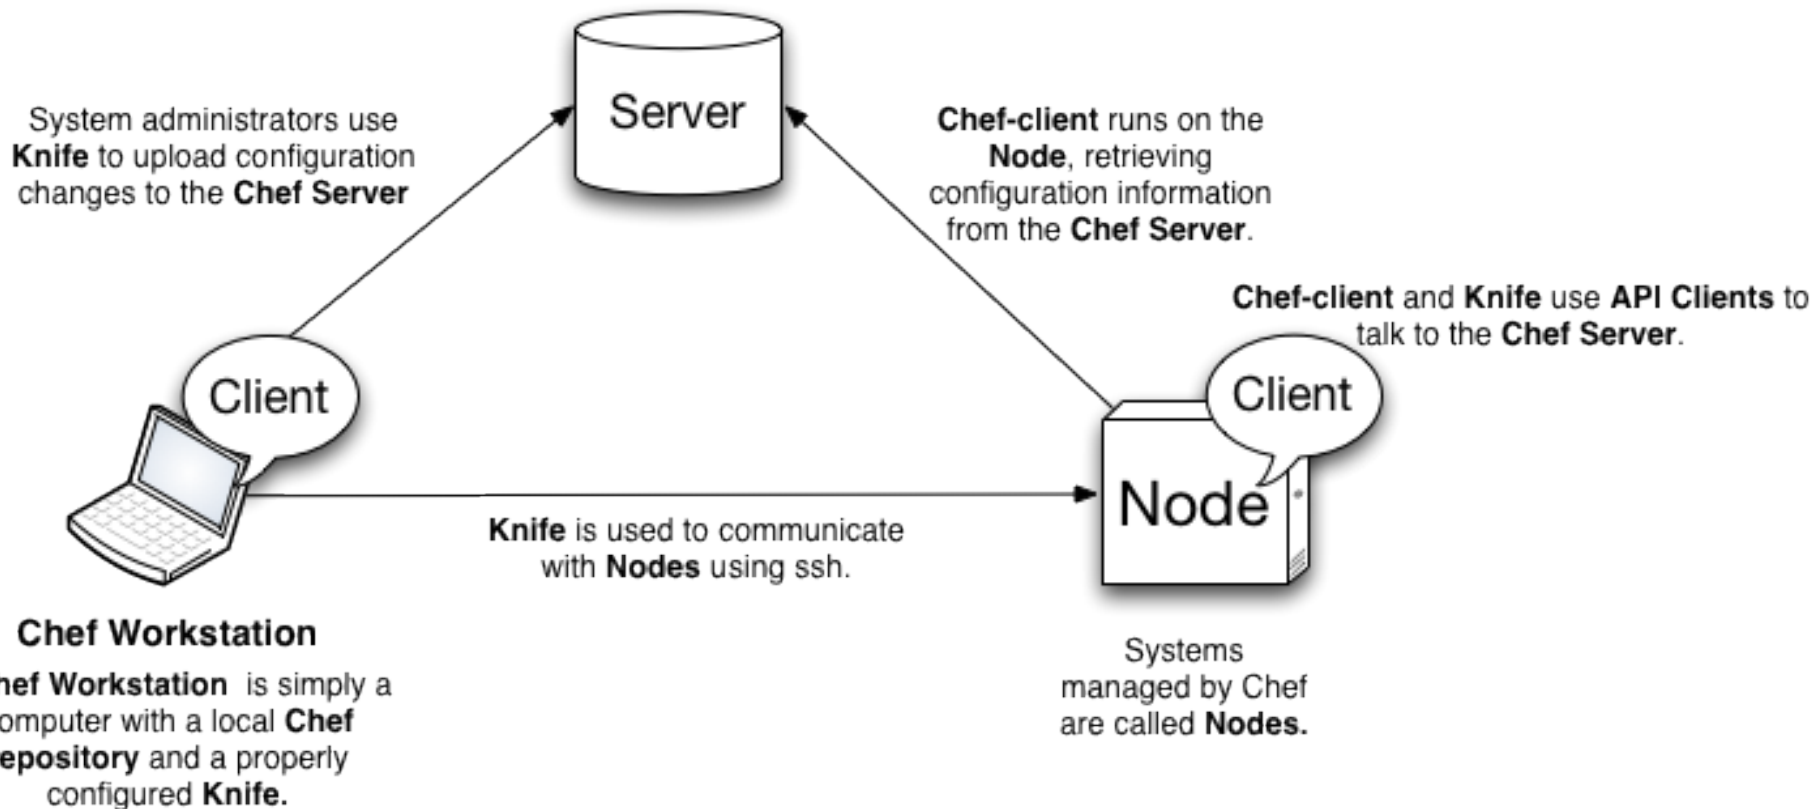
Built on top of Ruby

Chef Server Architecture



Architecture visualization

The **Chef Server** is the centralized store of your infrastructure's configuration.



Basic structure of Chef

Environments

Nodes

Cookbooks

Recipes

Files / Templates

Attributes

Data Bags

Search

What is what?

What is recipe?

What is cookbook?

What is resource?

What is node?

What is chef-client?

What is knife?

What is data bag?

What is template?

Modeling your infrastructure

Nodes

A **node** is a host that runs the Chef client. The primary features of a node, from Chef's point of view, are its **attributes** and its **run list**. Nodes are the thing that Recipes and Roles are applied to.

Roles

A **role** means grouping similar features of similar nodes.

Run list

A list of recipes that a node will run.

Configuring Nodes

Cookbooks

A **cookbook** is a collection of recipe, resource definition, attribute, library, cookbook file and template files that chef uses to configure a system. Cookbooks are typically grouped around configuring a single package or service.

The *MySQL* cookbook, for example, contains recipes for both client and server.

```
giedrius@lithuania:~/Projects/chef (master) $ tree -a cookbooks/mysql
```

```
cookbooks/mysql
├── README.rdoc
├── attributes
│   └── server.rb
├── metadata.rb
├── recipes
│   ├── client.rb
│   ├── default.rb
│   └── server.rb
└── templates
    ├── default
    │   ├── debian.cnf.erb
    │   ├── grants.sql.erb
    │   ├── my.cnf.erb
    │   ├── mysql-server.seed.erb
    │   └── slow_log_rotate.erb
```

```
4 directories, 11 files
```

Recipes

Recipes are the files where you write your resources (code).

```
package "nginx"

directory node[:nginx][:log_dir] do
  mode 0755
  owner node[:nginx][:user]
  action :create
end

%w{ngxsite nxdissite}.each do |nxscript|
  template "/usr/sbin/#{nxscript}" do
    source "#{nxscript}.erb"
    mode 0755
    owner "root"
    group "root"
  end
end
```

Another example

```
chef/cookbooks/git/recipes/default.rb
```

```
package "git-core"  
# apt-get install git-core  
# yum install git-core  
# etc..
```


Another example

```
directory "/home/new_folder" do
  mode 0755
  owner "someuser"
  group "www"
  action :create
end
```

Metadata

Cookbooks often rely on other cookbooks for pre-requisite functionality. In order for the server to know which cookbooks to ship to a client, a cookbook that depends on another one needs to express that dependency somewhere. That "somewhere" is in cookbook metadata.

Resources

A resource is usually a cross platform abstraction of the thing you're configuring on the host.

Chef's resources are mostly just containers for data, with some basic validation functionality.

Resources

Have a type

Have a name

Have parameters

Take action to put the resource in the declared state

Type

```
package "apache2" do
  version "2.2.11-2ubuntu2.6"
  action :install
end
```

```
template "/etc/apache2/apache2.conf" do
  source "apache2.conf.erb"
  owner "root"
  group "root"
  mode 0644
  action :create
end
```

Name

```
package "apache2" do
  version "2.2.11-2ubuntu2.6"
  action :install
end

template "/etc/apache2/apache2.conf" do
  source "apache2.conf.erb"
  owner "root"
  group "root"
  mode 0644
  action :create
end
```

Parameters

```
package "apache2" do
  version "2.2.11-2ubuntu2.6"
  action :install
end

template "/etc/apache2/apache2.conf" do
  source "apache2.conf.erb"
  owner "root"
  group "root"
  mode 0644
  action :create
end
```

Action

```
package "apache2" do
  version "2.2.11-2ubuntu2.6"
  action :install
end

template "/etc/apache2/apache2.conf" do
  source "apache2.conf.erb"
  owner "root"
  group "root"
  mode 0644
  action :create
end
```


Providers

The provider is the platform-specific implementation of the thing a resource abstracts.

On Red Hat or CentOS - yum

Debian and Ubuntu - apt package manager will be used

Search

Search is built by the Chef Server, and allow you to query arbitrary data about your infrastructure

Data Bags

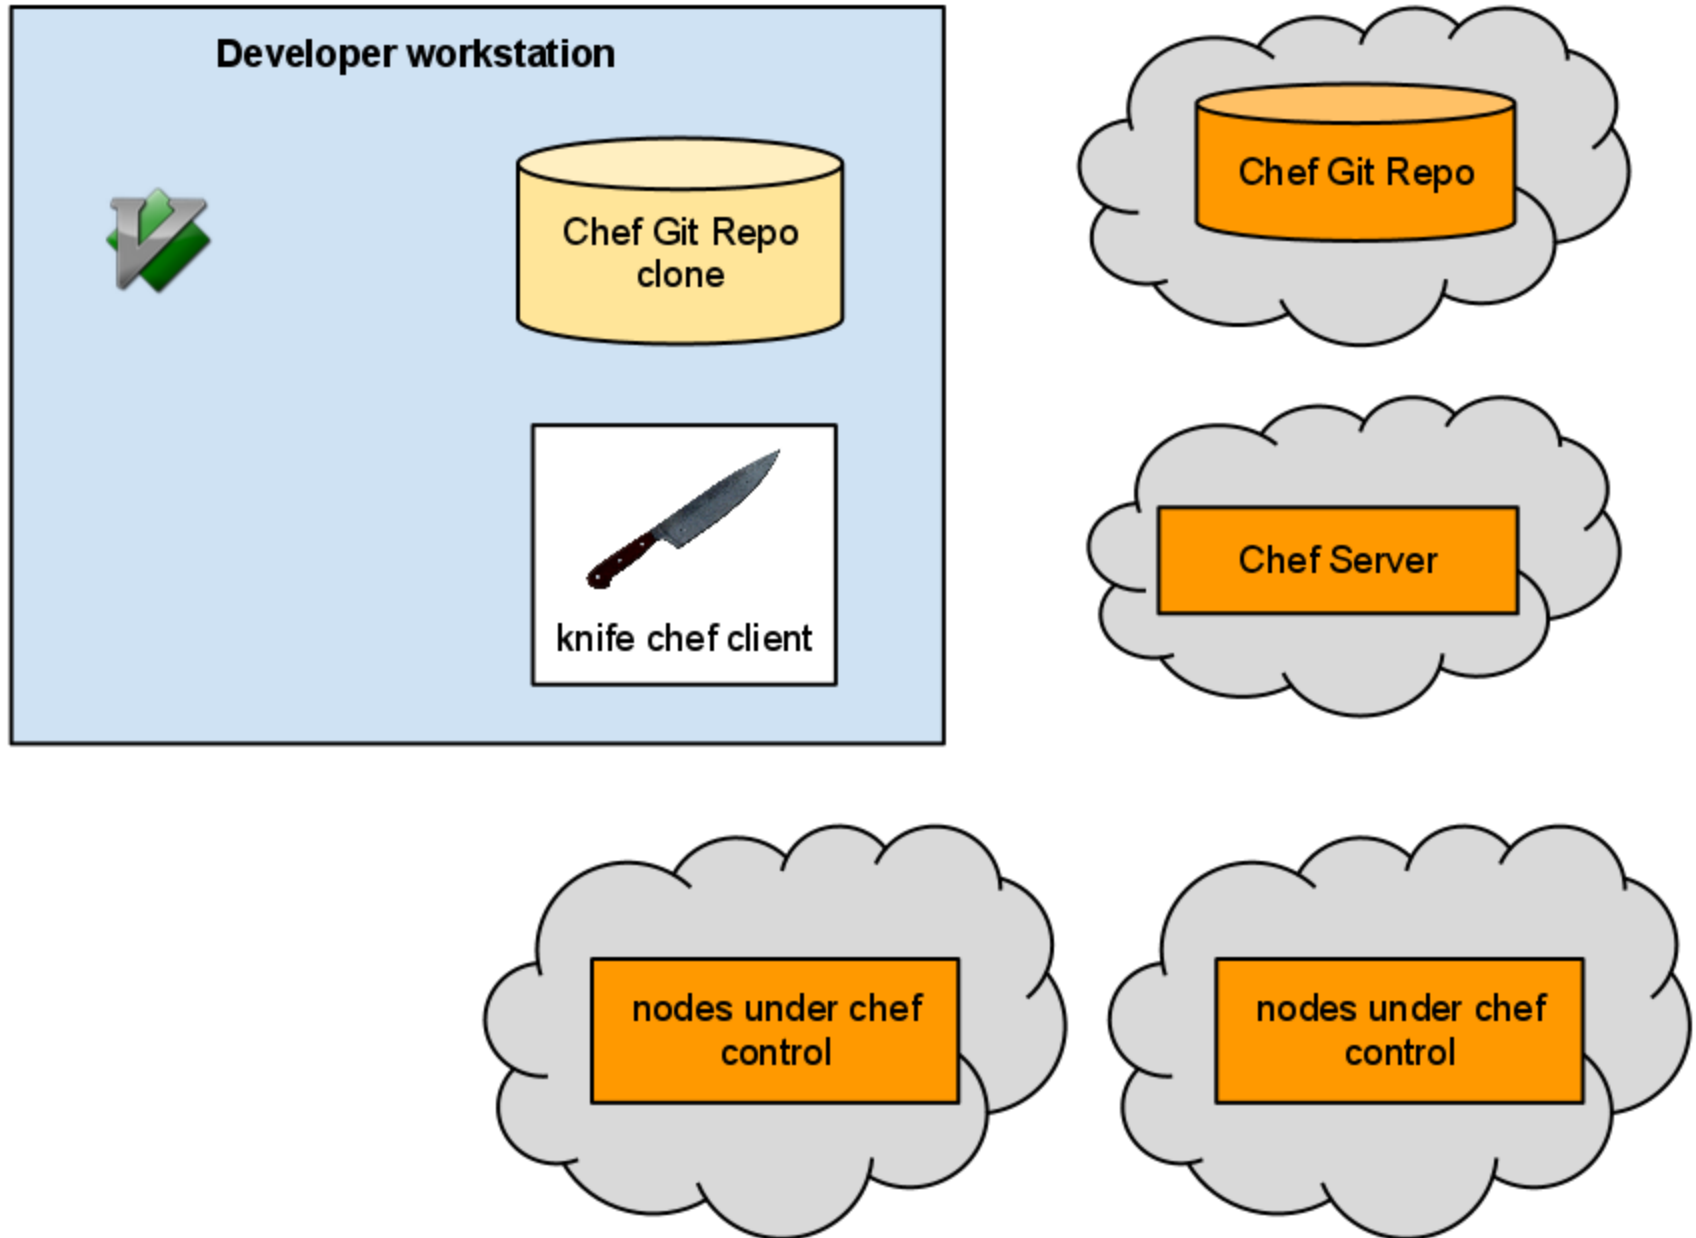
Data Bags store nested key-value data on the chef server. Data Bag data are searchable, and can also be loaded directly by name in a recipe. Data Bags are global for your chef-server installation-you can think of them as attributes for your whole infrastructure.

Environments

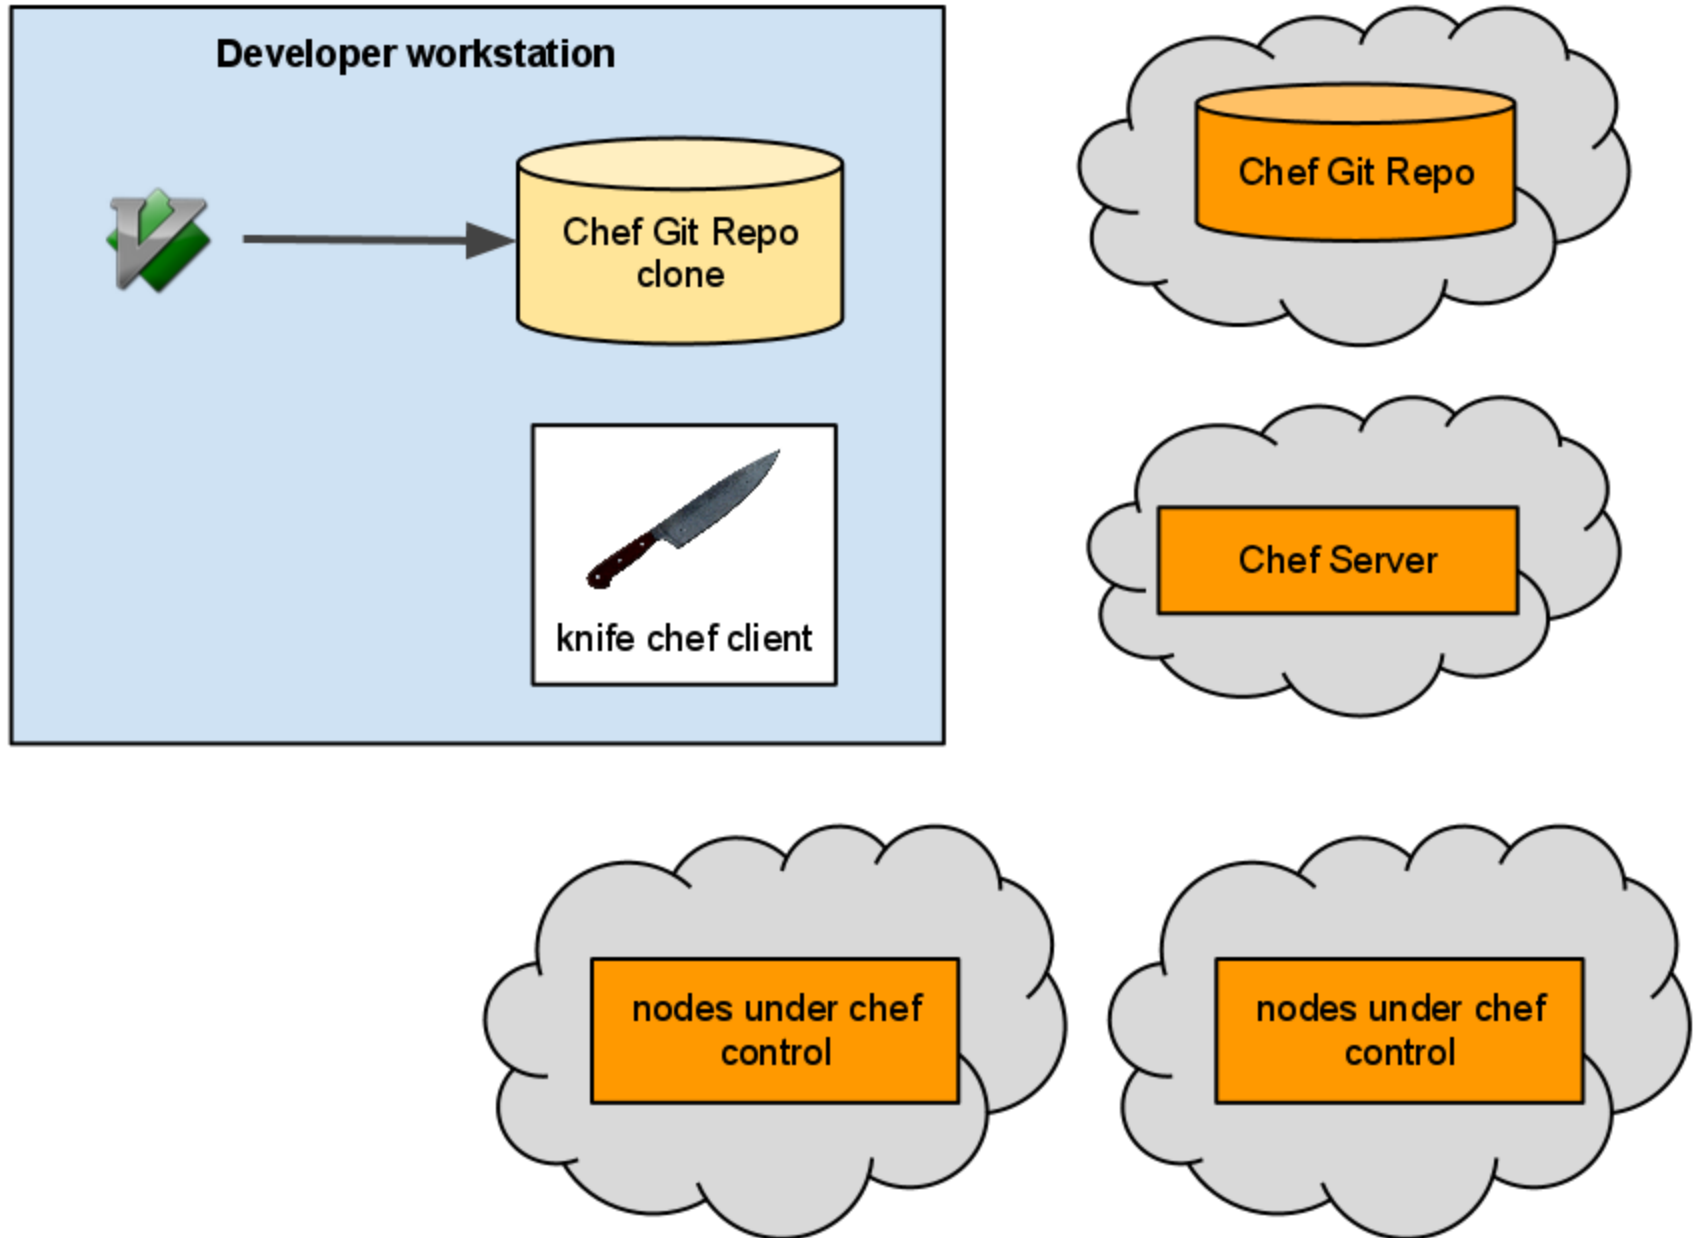
Provides a mechanism for managing different architectural segmented spaces such as production, staging, development, and testing, etc with one Chef setup.

Chef Workflow

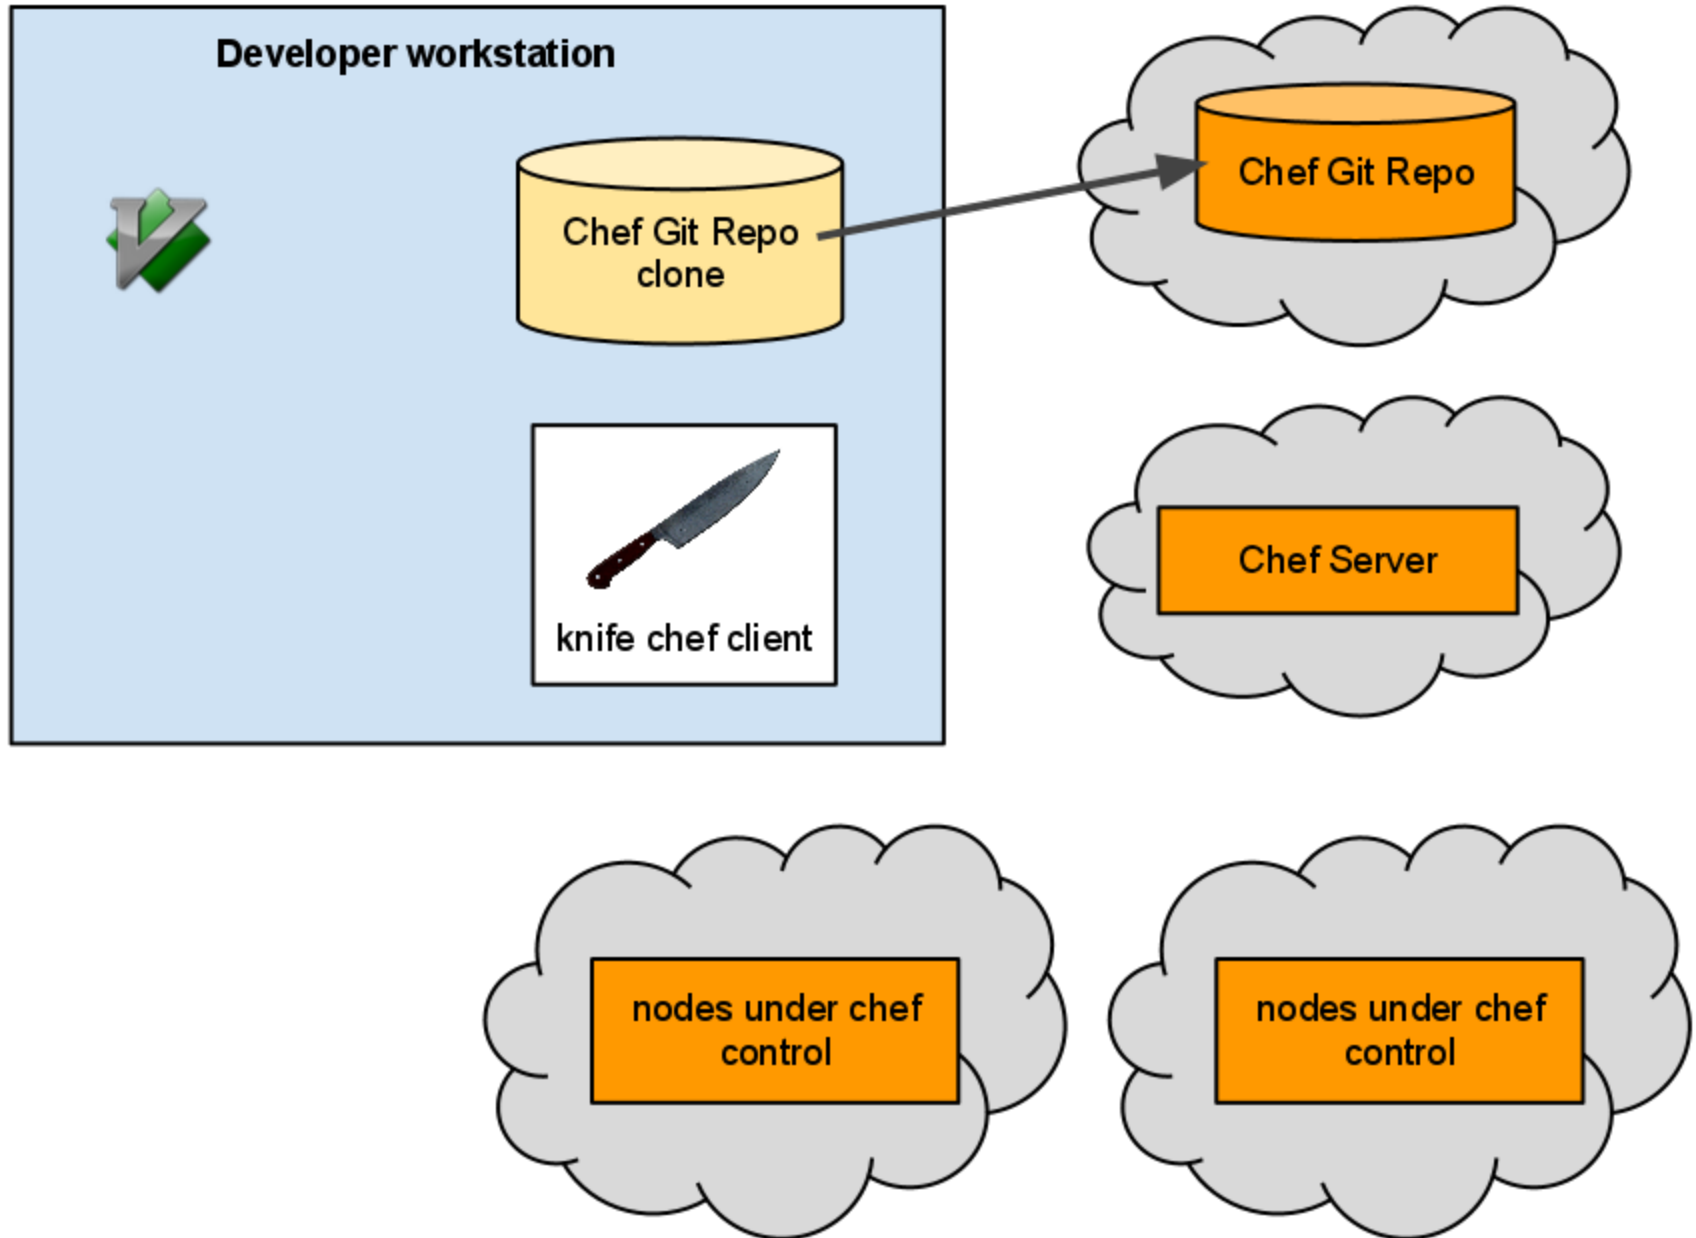
Everyday Chef Workflow for developers



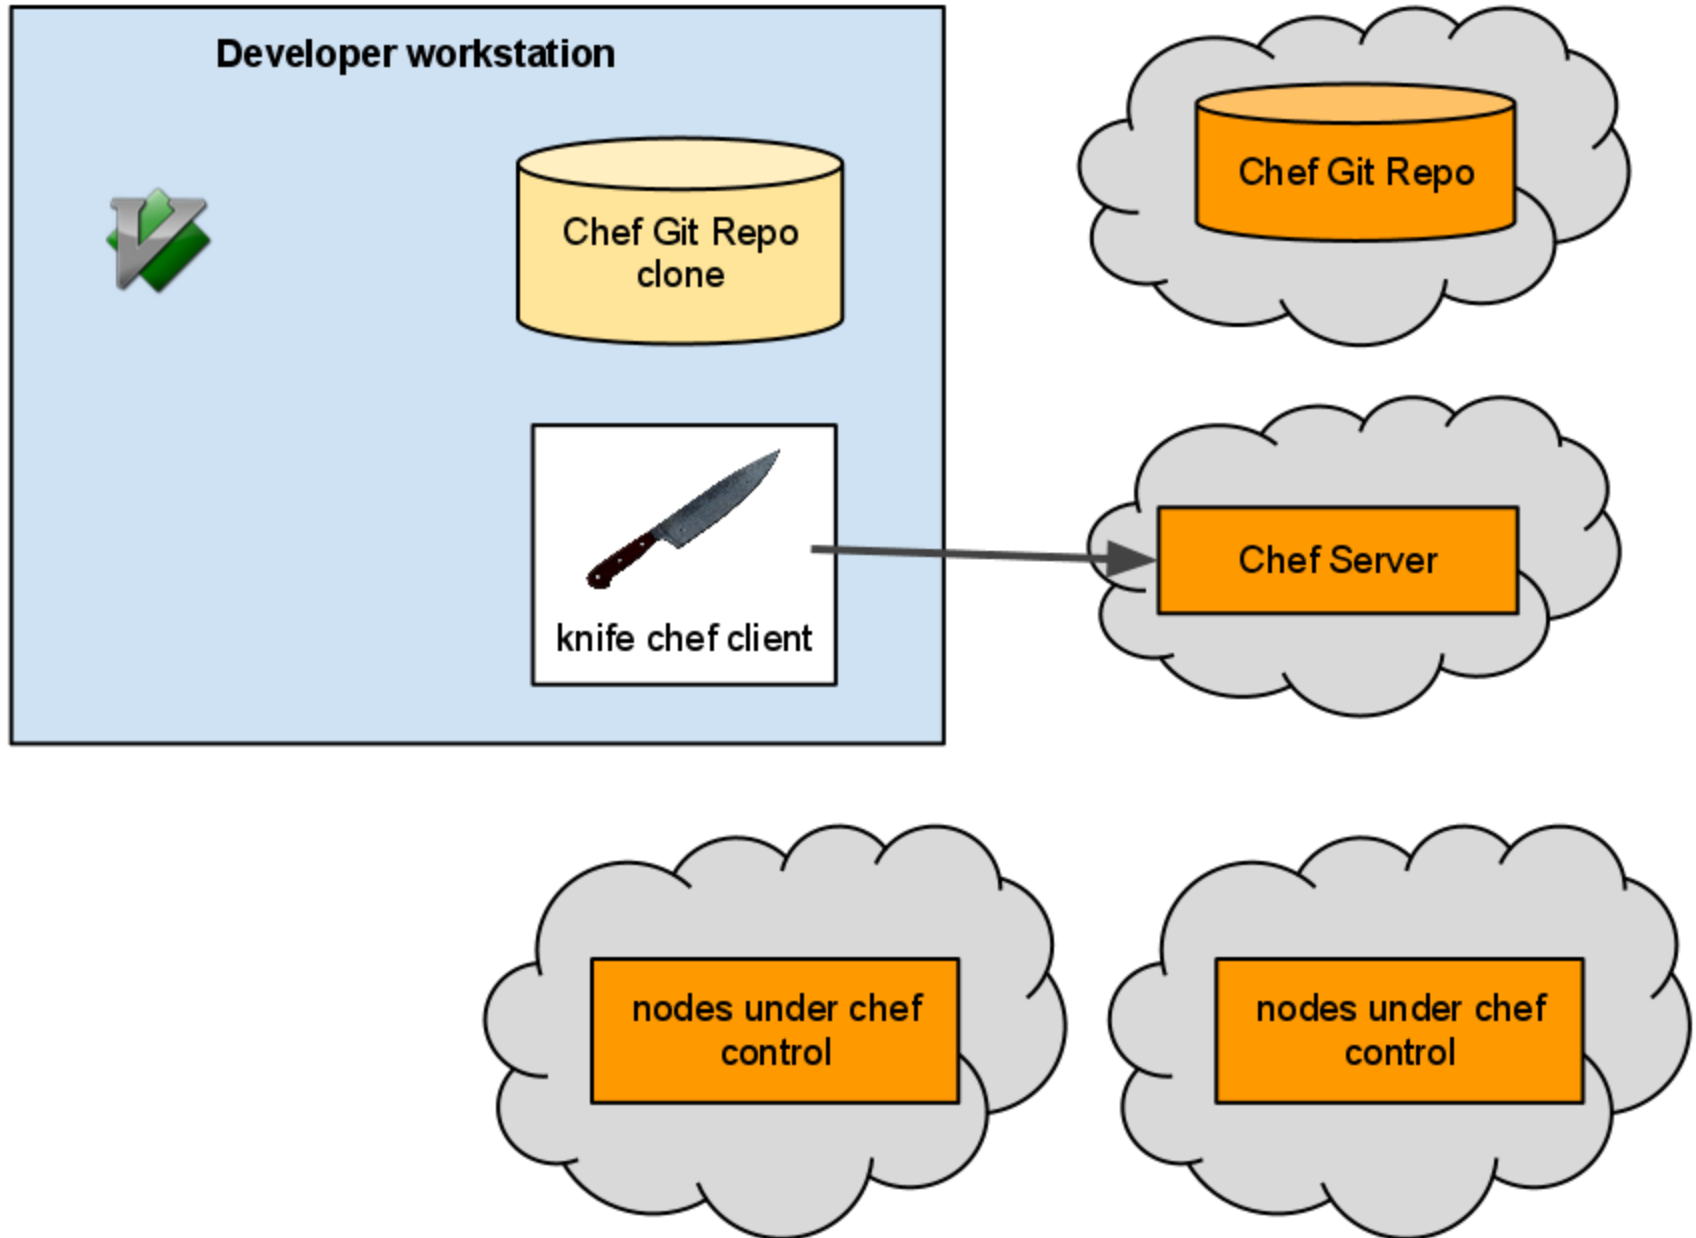
Developer makes changes to Chef recipes, roles, etc.



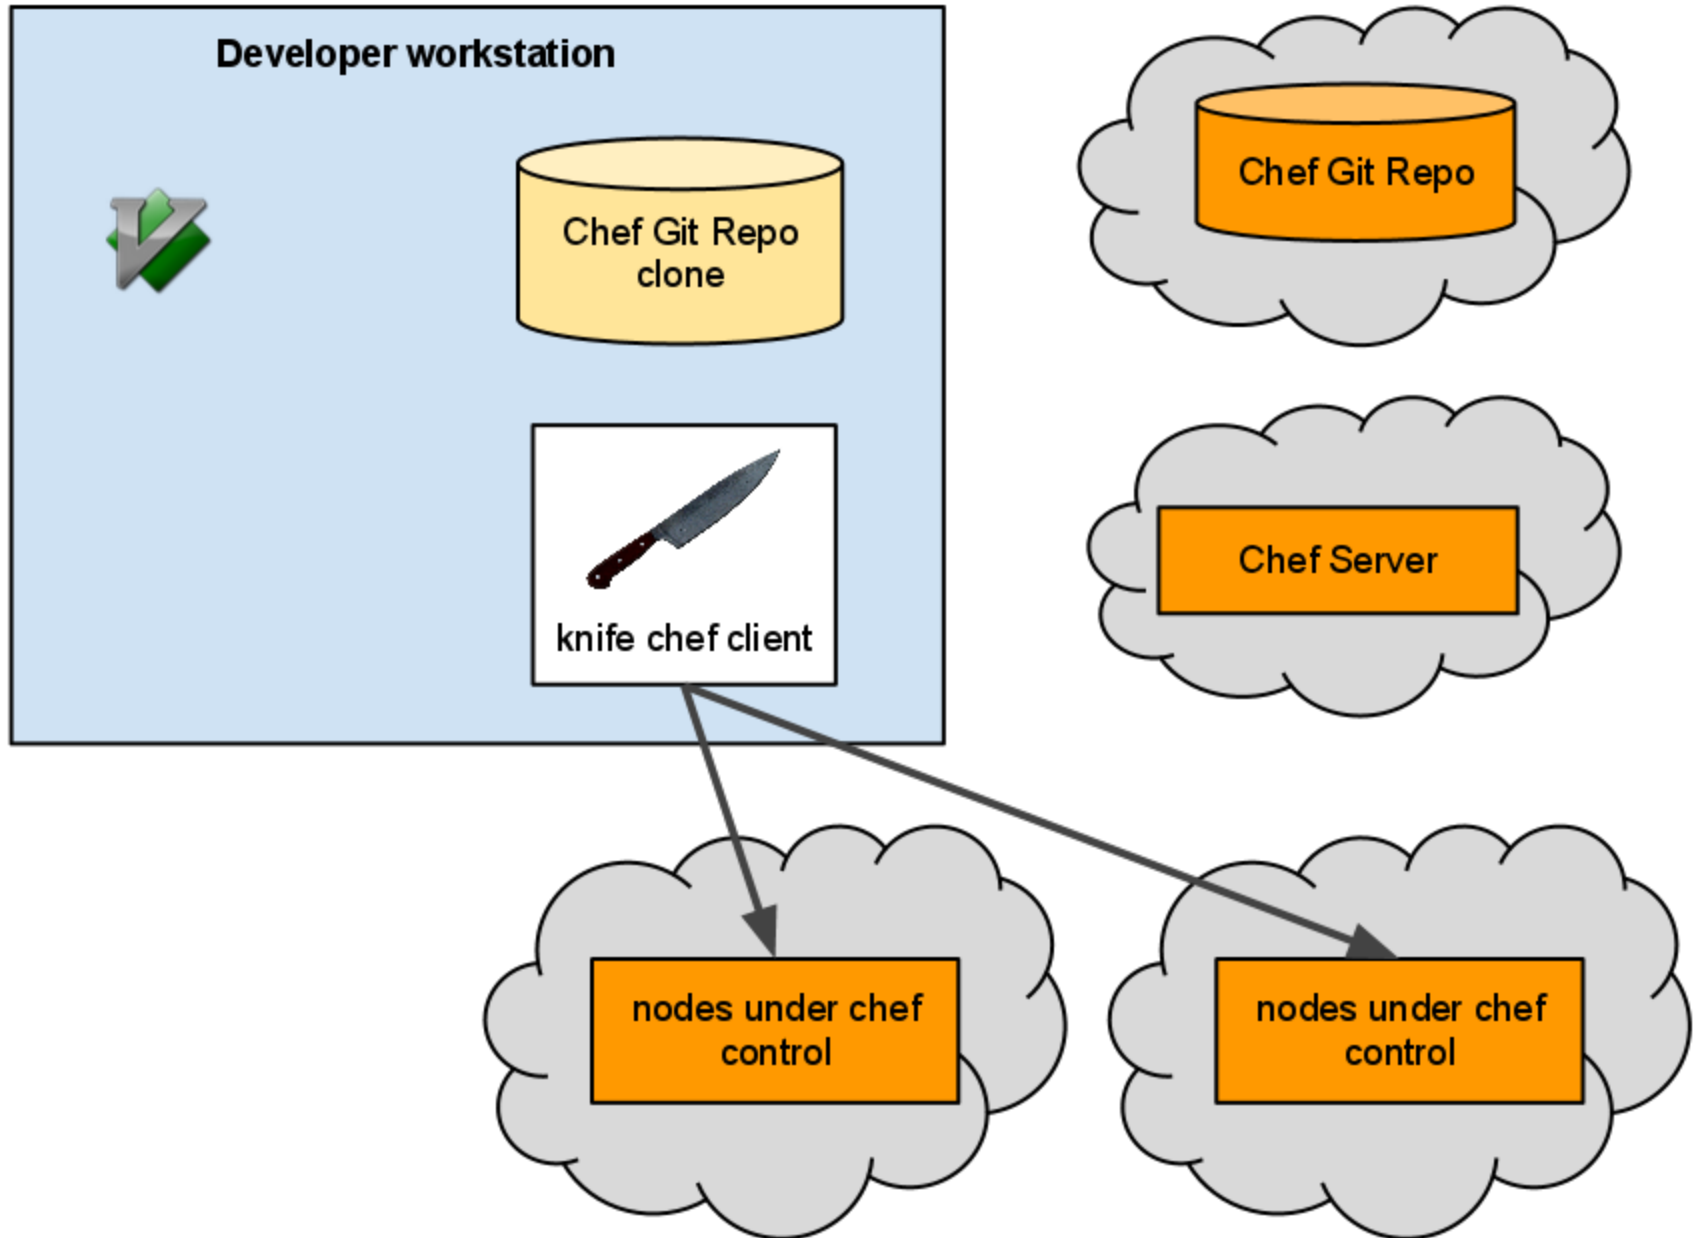
Developer pushes changes to origin Chef repository



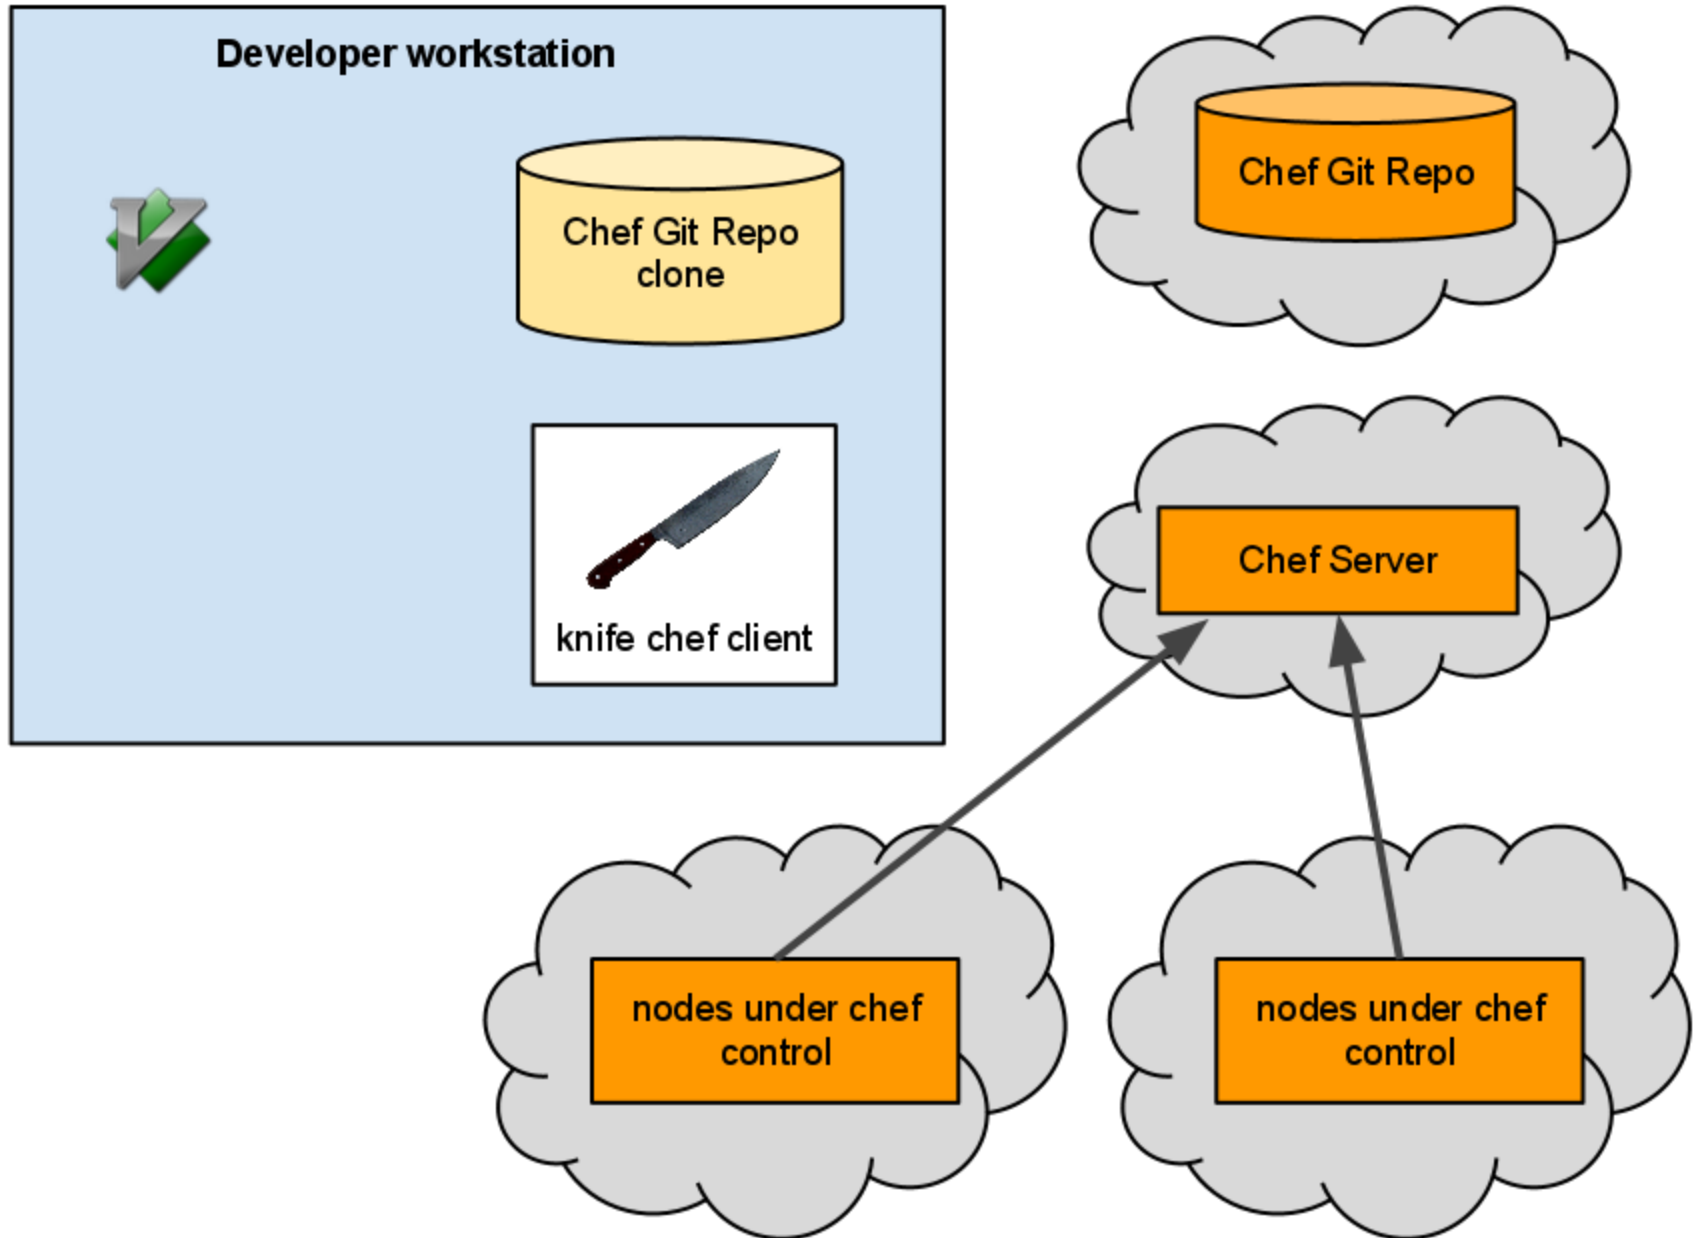
Developer uses Knife to push new code to the Chef Server



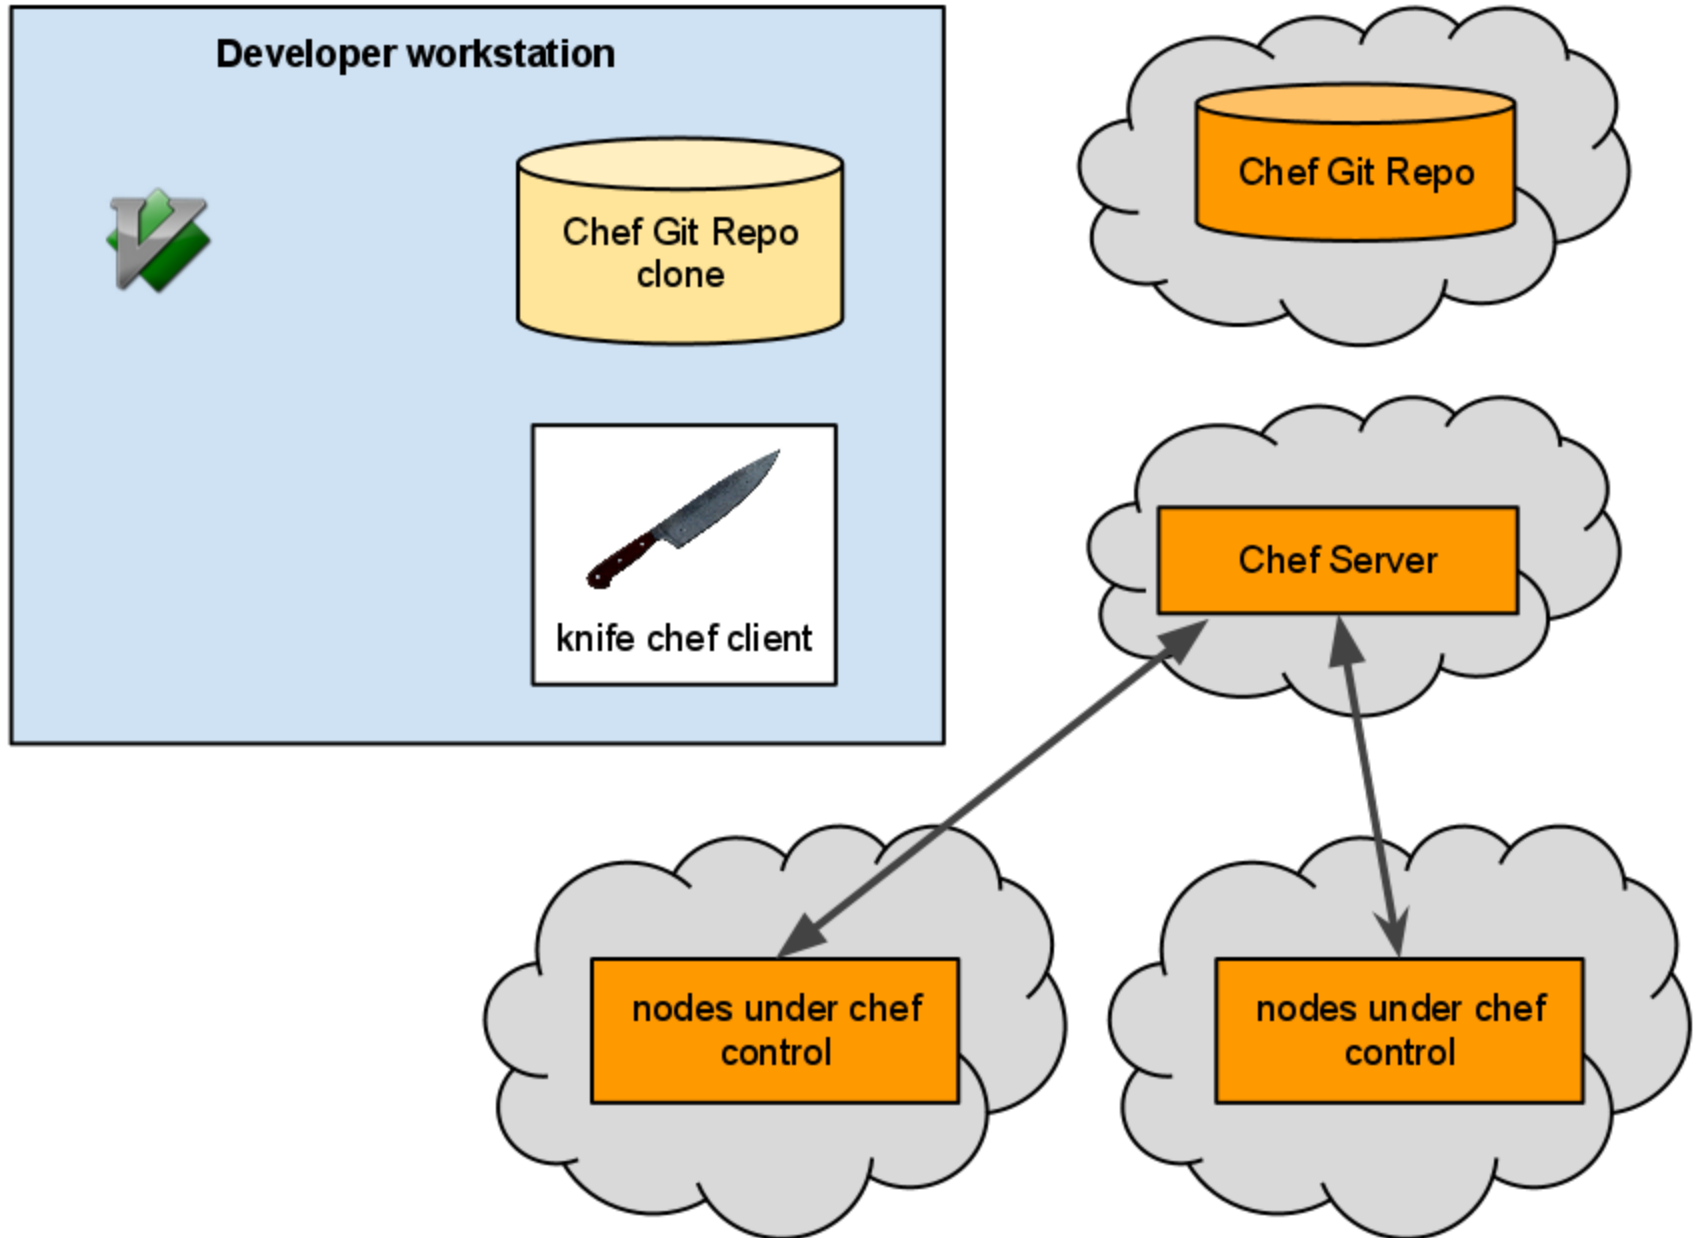
Developer uses Knife to tell Chef-clients to update themselves



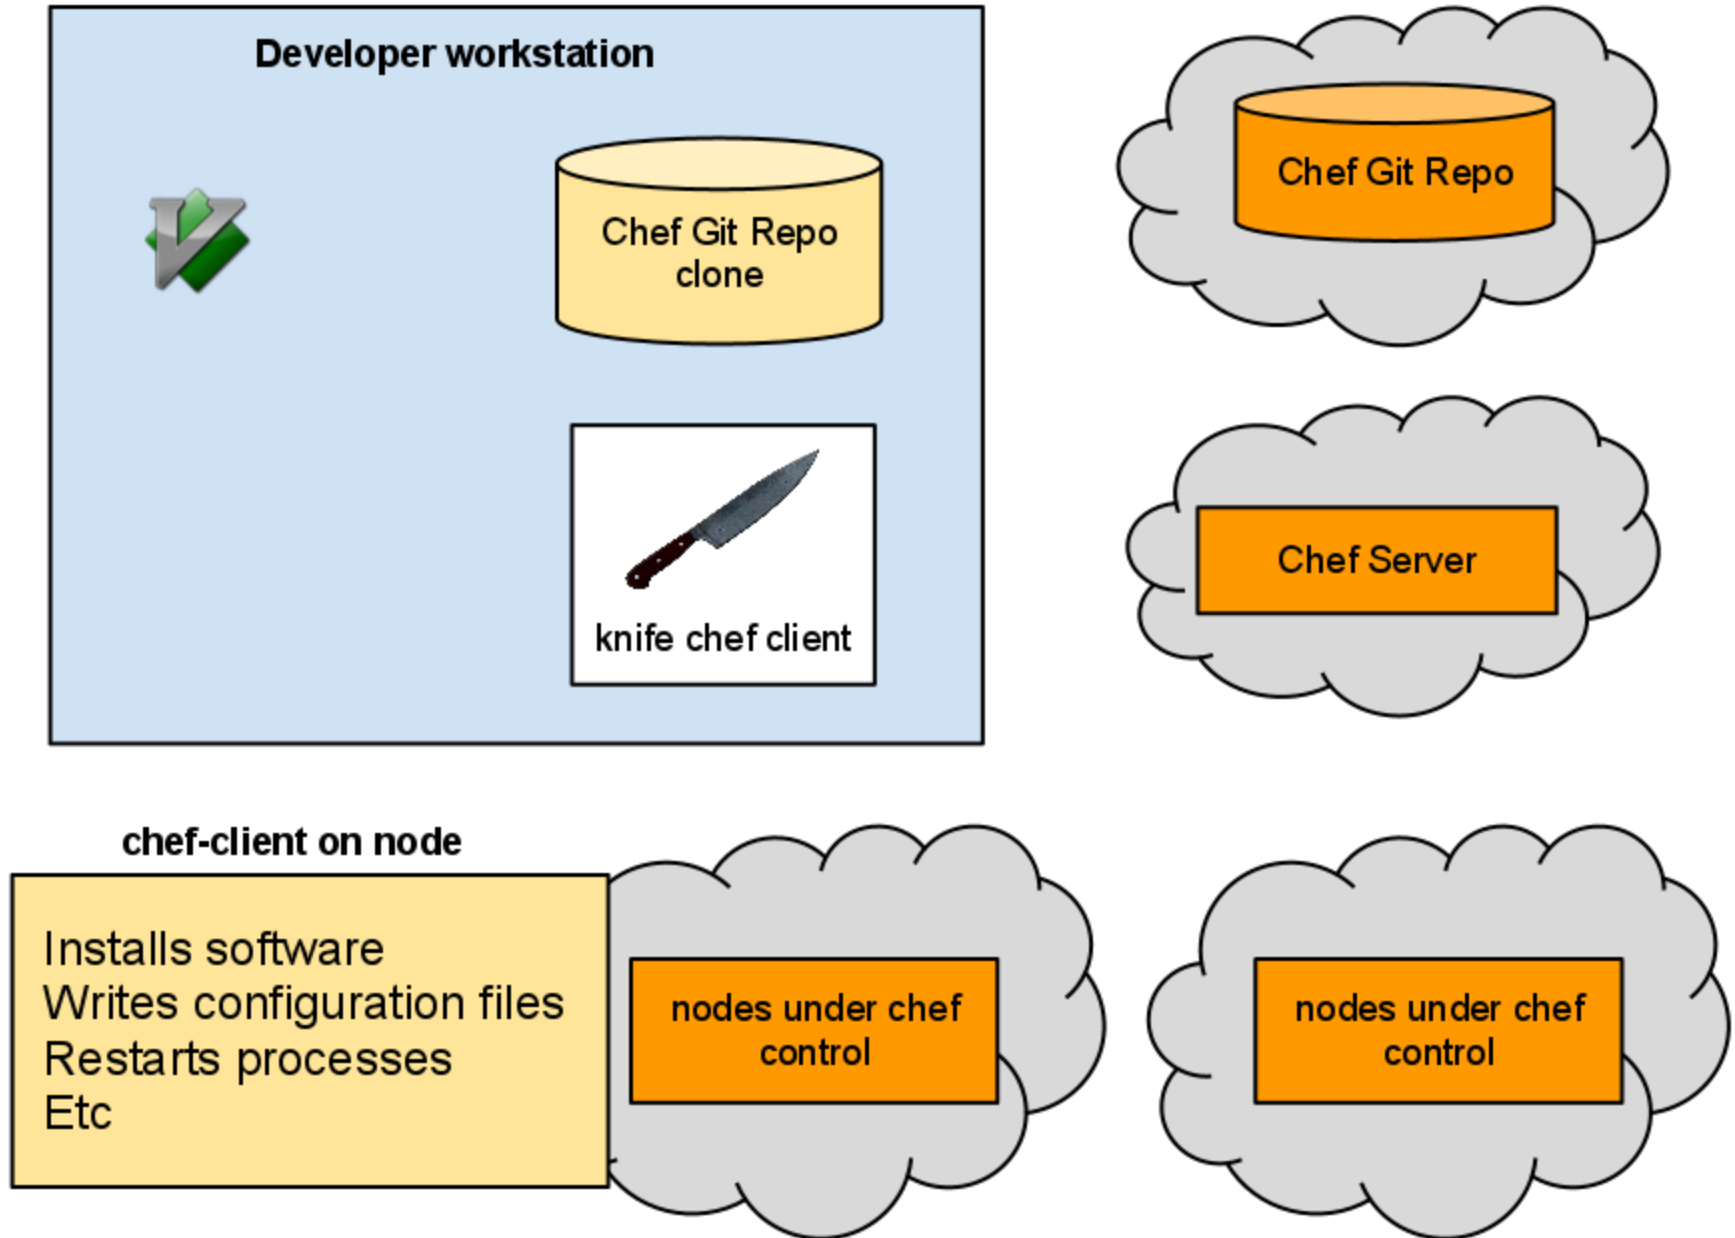
Chef-client on Chef nodes contact the Chef Server



Chef-client pulls latest code from the Chef Server



Chef-client on the node updates the system



Practical example

Update authorized_keys

```
$ cat cookbooks/bootstrap/files/default/authorized_keys
```

```
giedrius@lithuania:~/Projects/chef (master) $ cat cookbooks/bootstrap/files/default/authorized_keys  
ssh-rsa dev1-public-key  
ssh-rsa dev2-public-key  
ssh-rsa dev3-public-key
```

```
$ echo 'ssh-rsa dev4-public-key' >> !$
```

Committing and pushing

```
$ git diff
```

```
giedrius@lithuania:~/Projects/chef (master) $ git diff
diff --git a/cookbooks/bootstrap/files/default/authorized_keys b/cookbooks/bootstrap/files/default/authorized_keys
index 91cc004..8481dc6 100644
--- a/cookbooks/bootstrap/files/default/authorized_keys
+++ b/cookbooks/bootstrap/files/default/authorized_keys
@@ -1,3 +1,4 @@
 ssh-rsa dev1-public-key
 ssh-rsa dev2-public-key
 ssh-rsa dev3-public-key
+ssh-rsa dev4-public-key
_
```

```
git add .
git commit -m 'adding new public key'
git push
```


Updating Chef Server

```
bundle exec knife cookbook upload bootstrap
```

```
cap configure:all
```

or

```
ssh root@127.0.0.1
```

```
$ chef-client
```

How Chef helped us?

We don't care that much about

Infrastructure changes all the time

Scalability, because..

- adding new nodes is painless
- it's fast (takes minutes or hours, not days or weeks)
- there is no need in buying more machines with every new website

Clear separation between servers

with roles:

Frontend

Backend

Application

Memcached

Database

other..

Chef requirements

Chef-client is supported on the following platforms

- Ubuntu (10.04, 10.10, 11.04, 11.10)
- Debian (5.0, 6.0)
- RHEL & CentOS (5.x, 6.x)
- Fedora 10+
- Mac OS X (10.4, 10.5, 10.6, 10.7)
- Windows 7
- Windows Server 2003 R2, 2008 R2

Ruby

Ruby 1.8.7, 1.9.1 or 1.9.2 with SSL bindings is required.

RubyGems

Version 1.3.7 or greater. On Ubuntu and Debian Rubygems should be installed from source

Git?

Parts I didn't cover

- Setting up and running chef-server / chef-client
- Chef
- Knife plugins
- Chef + capistrano
- All the rest Chef goodies

Chef Alternatives

- . Puppet
- . Sprinkle
- . Rubber
- . Sunzi

Great resources

<https://github.com/opscode/cookbooks>

<http://wiki.opscode.com/display/chef/Home>

<http://railscasts.com/episodes/339-chef-solo-basics>

<https://github.com/ctshryock/capistrano-chef>

Thank You

tripsta®


travelplanet24®
τόσο απλή, τόσο φθηνά